

# COMP 330/543: Relational Databases 1

Luis Guzman

Sinan Kockara

Chris Jermaine

Rice University

# What is a Database?

A collection of data

Plus, a set of programs for managing that data

# Back in the Day...

The dominant data model was the network or navigational model (60's and 70's)

Data were a set of records with pointers between them

Much DB code was written in COBOL

Big problem was lack of physical data independence

- Code was written for specific storage model
  - Want to change storage? Modify your code
  - Want to index your data? Modify your code
  - Led to very little flexibility
- ▷ Your code locked you into a physical database design!

# Some People Realized This Was a Problem

By 1970, EF Codd (IBM) was looking at the so-called relational model

- Landmark 1970 paper, “A relational model of data for large shared data banks”
- Led to the 1981 Turing Award
  - ▷ Highest honor a computer scientist receives
  - ▷ Analogous to a Nobel Prize

Idea: data stored in “relations”

- A relation is a table of tuples or records
- Attributes of a tuple have no sub-structure (are atomic)

No pointers!

# Querying in the Relational Model

Querying is done via a “relational calculus”

Declarative

- You give a mathematical description of the tuples you want
- System figures out how to get those for you

Why is this good?

# Querying in the Relational Model

Querying is done via a “relational calculus”

Declarative

- You give a mathematical description of the tuples you want
- System figures out how to get those for you

Why is this good?

- Data independence!
- Your code has no data access specs
- So can change physical org, no code re-writes

# Relation Schema

All data are stored in tables, or relations

A relation schema consists of:

- A relation name (e.g., LIKES)
- A set of (attribute\_name, attribute\_type) pairs
  - ▷ Each pair is referred to as an “attribute”
  - ▷ Or sometimes as a “column”
- Usually denoted using LIKES (DRINKER string, BEER string)
- Or simply LIKES (DRINKER, BEER)

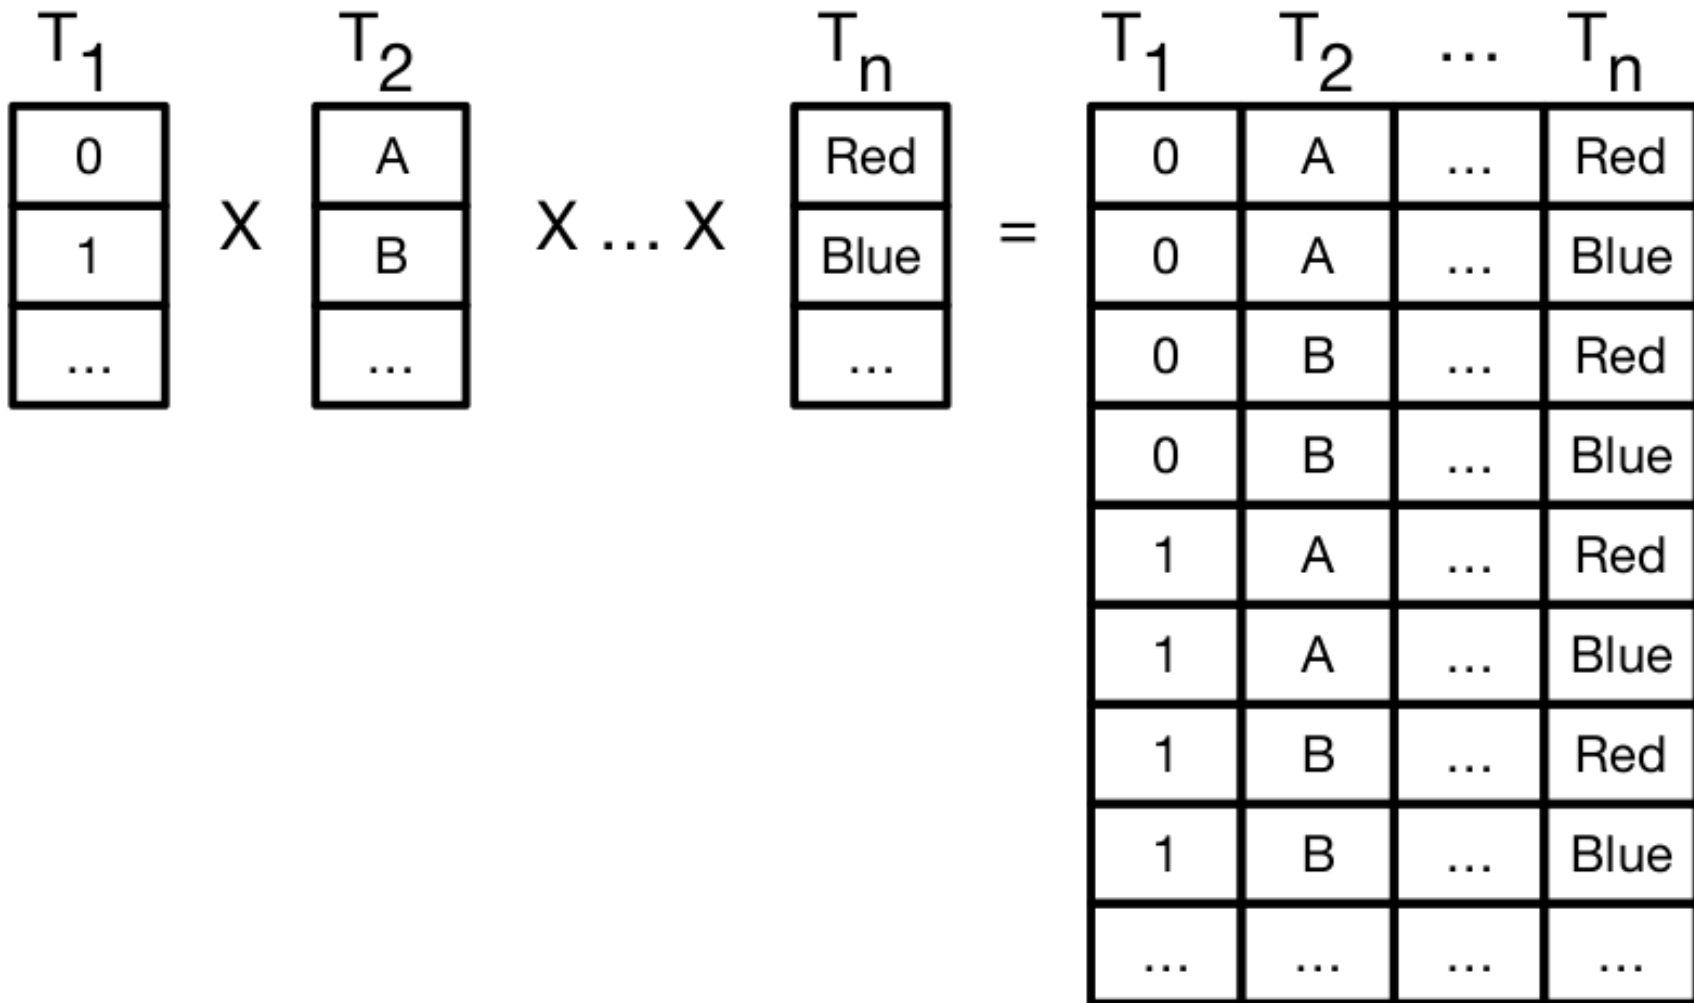
# A Relation

A relation schema defines a set of sets

- Specifically, if  $T_1, T_2, \dots, T_n$  are the  $n$  attribute types
- Where each  $T_i$  is a set of possible values
  - ▷ Ex: string is all finite-length character strings
  - ▷ Ex: integer is all numbers from  $-2^{31}$  to  $2^{31} - 1$
- Then a realization of the schema (aka a “relation”) is a subset of
  - ▷  $T_1 \times T_2 \times \dots \times T_n$
  - ▷ where  $\times$  is the Cartesian product operator



# Attribute Types Forming a Relation



## A Relation (continued)

So for the relation schema LIKES (DRINKER string, BEER string)

A corresponding relation might be

$$\{(\text{“Luis”}, \text{“Modelo”}),$$
$$(\text{“Sinan”}, \text{“PBR”})\}$$

This is also referred to as a “table”

The entries in the relation are referred to as

- “rows”
- “tuples”
- “records”

# Keys

In the relational model, given  $R(A_1, A_2, \dots, A_n)$

A set of attributes  $K = \{K_1, \dots, K_m\}$  is a KEY of  $R$  if:

- For any valid realization  $R'$  of  $R$ ...
- For all  $t_1, t_2$  in  $R'$ ...
- If  $t_1[K_1] = t_2[K_1]$  and  $t_1[K_2] = t_2[K_2]$  and ...  $t_1[K_m] = t_2[K_m]$ ...
- Then it must be the case that  $t_1 = t_2$

Note: every relation schema MUST have a key... why?

# Keys

What is a key for

STUDENT (NETID, FNAME, LNAME, AGE, COLLEGE)?

What is a key for

LIKES (DRINKER, BEER)?

# Keys (continued)

A relation schema can have many keys

Those that are minimal are CANDIDATE KEYS

- “Minimal” means no subset is a key


One is designated as the PRIMARY KEY denoted with an underline

- STUDENT (NETID, FNAME, LNAME, AGE, COLLEGE)

Surrogate Key

- No real-world meaning, added to simplify primary key

# Keys (continued)



# Connecting Relations

The relational model does not have pointers

Why? Two reasons:

- Not nice mathematically
  - ▷ Mathematical elegance key goal in model design
- Implementation difficult
  - ▷ Move an object? All pointers are invalid!
  - ▷ Can have centralized look-up table
  - ▷ But expensive, complicated, plus problem still exists

# Connecting Relations

But we still need some notion of between-tuple references

- DRINKERS (DRN\_ID, FNAME, LNAME)
- LIKES (DRN\_ID, BEER)

Clearly, LIKES.DRN\_ID refers to DRINKERS.DRN\_ID

- Why not the other way around?

Accomplished via the idea of a FOREIGN KEY



# Solution: Foreign Keys

- DRINKERS (DRN\_ID, FNAME, LNAME)
- LIKES (DRN\_ID, BEER)

Given relation schemas  $R_1, R_2$

- We say a set of attributes  $K_1$  from  $R_1$  is a foreign key to a set of attributes  $K_2$  from  $R_2$  if...
- (1)  $K_2$  is a candidate key for  $R_2$ , and...
- (2) For any valid realizations  $R'_1, R'_2$  of  $R_1, R_2$ ...
- For each  $t_1 \in R'_1$ , it MUST be the case that there exists  $t_2 \in R'_2$  s.t...
- $t_1[K_{1,1}] = t_2[K_{2,1}]$  and  $t_1[K_{1,2}] = t_2[K_{2,2}]$  and ...  $t_1[K_{1,m}] = t_2[K_{2,m}]$

Intuitively, what does this mean?

# Foreign Keys (continued)

Intuitively what does it mean?

- The foreign key ( $k_1$ ) must be a set of attributes that uniquely identify a record in another table ( $R_2$ )
- The combination of attribute values present in every tuple of ( $R_1$ ) must also be present in ( $R_2$ )

Why is this a requirement?

# Foreign Keys (continued)

Intuitively what does it mean?

- The foreign key ( $k_1$ ) must be a set of attributes that uniquely identify a record in another table ( $R_2$ )
- The combination of attribute values present in every tuple of ( $R_1$ ) must also be present in ( $R_2$ )

Why is this a requirement?

- To prevent inconsistencies
- To match to a single target

RDBMS enforce these requirements via

- Cascading deletes
- Failed inserts

## Foreign Keys (continued)

Which one is  $R_1$  and  $R_2$ ?

DRINKERS

DRN_ID	FNAME	LNAME
lg67	Luis	Guzman
sk212	Sinan	Kockara

LIKES

DRN_ID	BEER
lg67	Modelo
sk212	PBR
lg67	Corona

## Foreign Keys (continued)

What happens here?

### DRINKERS

DRN_ID	FNAME	LNAME
lg67	Luis	Guzman
sk212	Sinan	Kockara

### LIKES

DRN_ID	BEER
lg67	Modelo
sk212	PBR
lg67	Corona
cmj4	Blue Moon

# Queries/Computations in the Relational Model

The original query language was the RELATIONAL CALCULUS

- Fully declarative programming language

next was the RELATIONAL ALGEBRA

- Imperative
- Define a set of operations over relations
- A RA program is then a sequence of those operations
- This is the “abstract machine” of RDBs

Today we use SQL

- Heavily influenced by RC
- Has aspects of RA

# Overview of Relational Calculus

RC is a variant on first-order logic

You say: “Give me all tuples  $t$  where  $P(t)$  holds”

$P(t)$  is a predicate in first-order logic

- A predicate is basically boolean function

# Predicates

First order logic allows predicates

- ▷ predicate: a function that evals to true/false
- ▷ “It’s raining on day  $X$ ” or *Raining*( $X$ )
- ▷ “It’s cloudy on day  $X$ ” or *Cloudy*( $X$ )

Can build more complicated preds using logical operations over them

- ▷ and ( $\wedge$ )
- ▷ or ( $\vee$ )
- ▷ not ( $\neg$ )
- ▷ implies ( $\rightarrow$ )
- ▷ iff ( $\leftrightarrow$ )



# Predicates (continued)

Example:  $Raining(X) \wedge Cloudy(X)$

Evals to true if:

- ▷ It is raining and cloudy on day  $X$

Example:  $Raining(X) \rightarrow Cloudy(X)$

Evals to true if either:

- ▷ It is not raining on day  $X$ , or
- ▷ It is raining and cloudy on day  $X$

Note the difference between them!

- ▷  $\rightarrow$  is like a logical “if-then”

# First Order Logic

Just predicates and logical ops?

▷ You've got predicate logic

But when you add quantification

▷  $\forall, \exists$

▷ You've got first order logic

# Universal Quantification

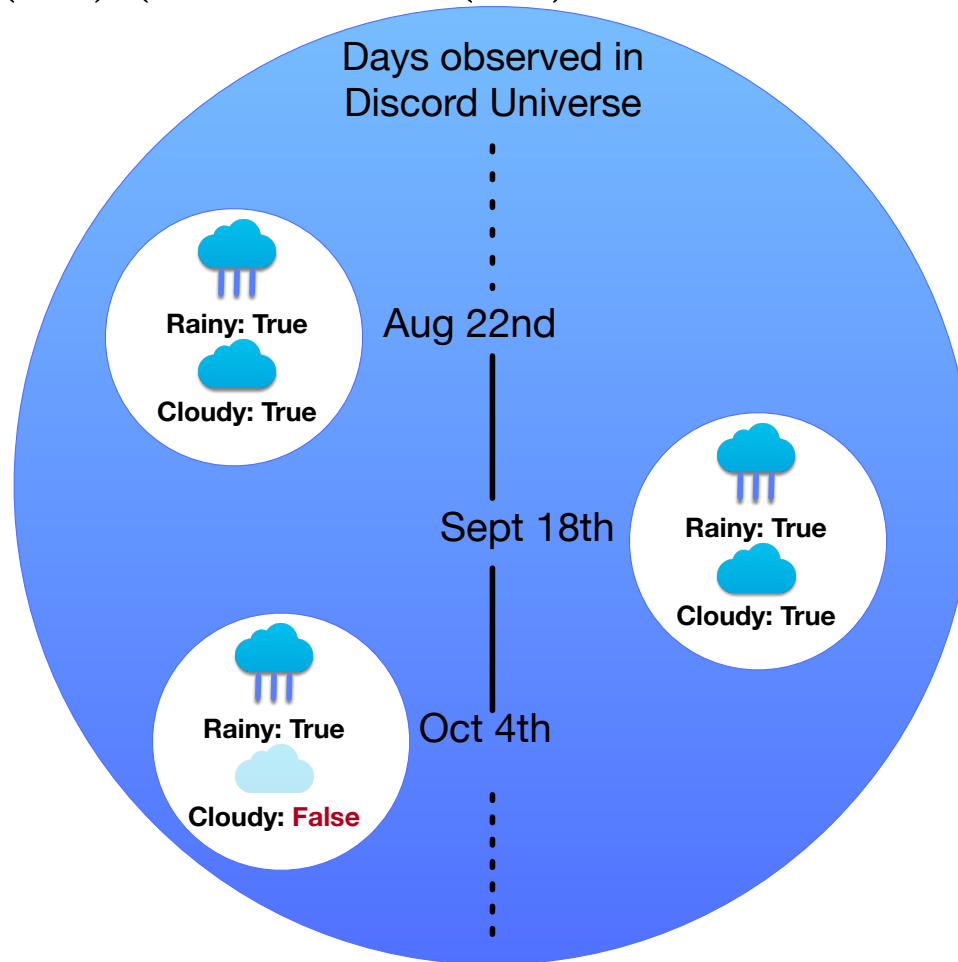
Asserts that a predicate is true all of the time

Example:

- ▷  $\forall(X)(\textit{Raining}(X) \rightarrow \textit{Cloudy}(X))$
- ▷ This is a zero-arg predicate (takes no params)
- ▷ Asserts that it only rains when it is cloudy
- ▷ Note: idea of universe of discourse is key!

# Universal Quantification: Example 1

$$\forall(X)(\textit{Raining}(X) \rightarrow \textit{Cloudy}(X))$$




# Universal Quantification

Asserts that a predicate is true all of the time

- ▷  $\forall(X)(\textit{Friends}(X, Y))$
- ▷ This is a predicate over  $Y$
- ▷ Eval's to true if the person  $Y$  is friends with everyone

# Universal Quantification: Example 2

$$\forall(X)(\textit{Friends}(X, Y))$$



# Existential Quantification

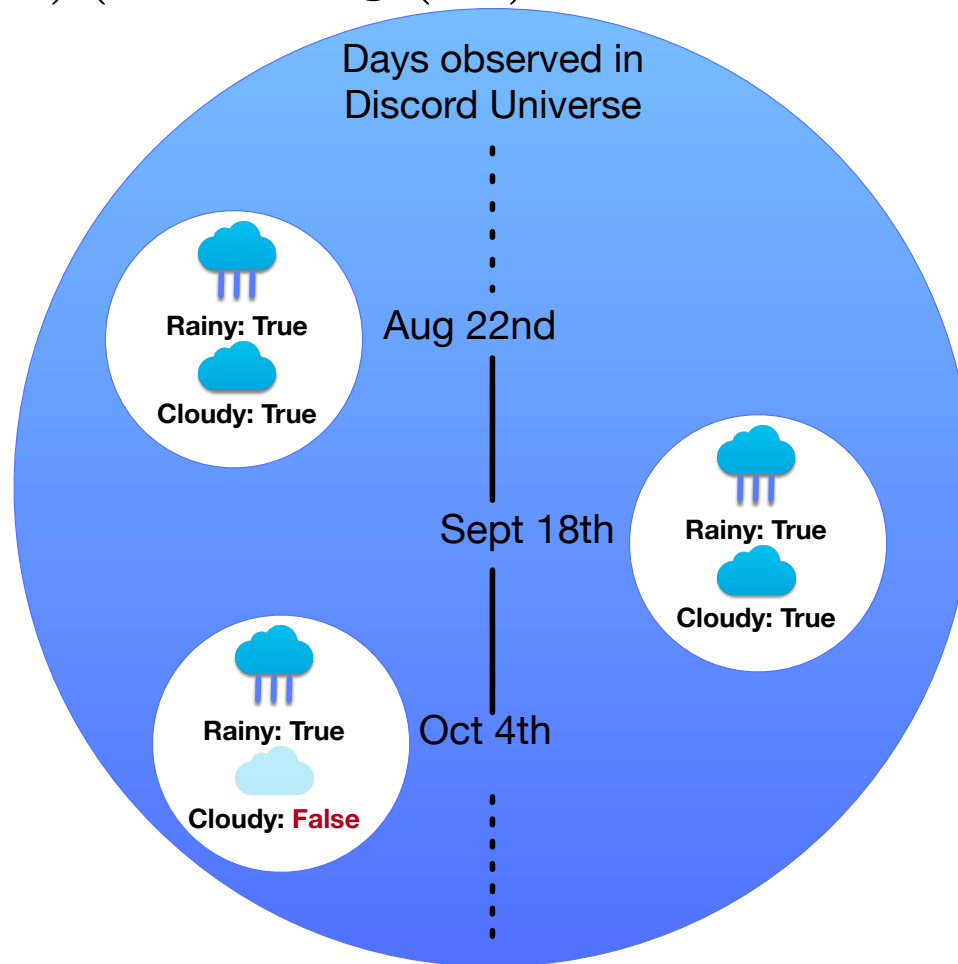
Asserts that a predicate can be satisfied

Example:

- ▷  $\exists(X)(\textit{Raining}(X) \wedge \text{not } \textit{Cloudy}(X))$
- ▷ Asserts that it can rain when it is not cloudy
- ▷ A.K.A a "sun shower"

# Existential Quantification: Example 1

$$\exists(X)(\textit{Raining}(X) \wedge \text{not } \textit{Cloudy}(X))$$





# Existential Quantification

Asserts that a predicate can be satisfied

Example:

- ▷  $\text{not } \exists(X, Y)(\text{Friends}(X, Y) \wedge \text{Friends}(X, Z) \wedge \text{Friends}(Y, Z))$
- ▷ This is a predicate over  $Z$
- ▷ Evals to true when?

# Existential Quantification

Asserts that a predicate can be satisfied

Example:

- ▷  $\text{not } \exists(X, Y)(\text{Friends}(X, Y) \wedge \text{Friends}(X, Z) \wedge \text{Friends}(Y, Z))$
- ▷ This is a predicate over  $Z$
- ▷ Eval's to true when? (If there is not a pair of friends who are both friendly with  $Z$ )

# Important Equivalence

$\forall(X)(P(X))$  is equivalent to... **not**  $\exists(X)(\text{not } P(X))$

- ▷ Ex: **not**  $\exists(X, Y)(\textit{Friends}(X, Y) \wedge \textit{Friends}(X, Z) \wedge \textit{Friends}(Y, Z))$
- ▷ Can be changed to:
- ▷  $\forall(X, Y)(\text{not } (\textit{Friends}(X, Y) \wedge \textit{Friends}(X, Z) \wedge \textit{Friends}(Y, Z)))$
- ▷ Or  $\forall(X, Y)(\text{not } \textit{Friends}(X, Y) \vee \text{not } \textit{Friends}(X, Z) \vee \text{not } \textit{Friends}(Y, Z))$

Which is easier?

- IMO: often easier to reason about  $\forall$
- With  $\exists$  tend to get into double and triple negatives
- Unfortunately, SQL does not have  $\forall$

# Questions?