

# COMP 330/543: SQL 1

Luis Guzman

Sinan Kockara

Chris Jermaine

Rice University

# SQL

De-facto standard DB programming language

- ▷ First proposed by IBM researchers in 1970's
- ▷ Oracle first to offer commercial version in 1979
- ▷ IBM soon after
- ▷ Donald D. Chamberlin Video

SQL is a H U G E language!!

- ▷ Current standard runs to 100s of pages
- ▷ Consists of a declarative DML
- ▷ And an imperative DML
- ▷ And a DDL

# Relational Calculus/Algebra vs. SQL

Duplicates are not automatically eliminated

Not all SQL implementations are compatible

- ▷ Support different set of operators
- ▷ Date and time syntax
- ▷ Comparison case sensitivity

SQL extends RC/RA

- ▷ Aggregate functions
- ▷ Schema modifications

# RA vs SQL Operators

RA Name	RA symbol	SQL equivalent
Projection	$\pi$	SELECT [L: attribute list]
Join	$\times \bowtie *$	FROM [R: Relation list]
Selection	$\sigma$	WHERE [C: Condition list]
Union	$\cup$	UNION or UNION ALL
Intersection	$\cap$	JOIN or EXISTS or IN
Difference	$-$	EXCEPT
Rename	$\rho$	AS
Assignment	$\leftarrow$	INTO

# Query Structure

We begin with the heart and soul of SQL: the declarative DML

```
SELECT <attribute list>  
FROM <tables>  
WHERE <conditions> (Optional)
```

```
SELECT DRINKER, BEER FROM LIKES
```

DRINKER	BEER
Chris	Double Trouble
Chris	Tout Suite
Luis	Blue Moon
Luis	Modelo

# Query Structure

We begin with the heart and soul of SQL: the declarative DML

```
SELECT <attribute list>  
FROM <tables>  
WHERE <conditions>
```

```
SELECT DRINKER, BEER FROM LIKES WHERE DRINKER = "Luis"
```

DRINKER	BEER
Luis	Blue Moon
Luis	Modelo

# Our First “Real” Query

LIKES (DRINKER, BEER)

FREQUENTS (DRINKER, BAR)

SERVES (BAR, BEER)

Who goes to a bar serving Sam Smith Taddy Porter? (‘SSTP’)

**SELECT**

**FROM**

**WHERE**

# Our First “Real” Query

LIKES (DRINKER, BEER)

FREQUENTS (DRINKER, BAR)

SERVES (BAR, BEER)

Who goes to a bar serving Sam Smith Taddy Porter? (‘SSTP’)

```
SELECT f.DRINKER  
FROM FREQUENTS AS f, SERVES AS s  
WHERE f.BAR = s.BAR AND s.BEER = "SSTP"
```

Are we missing anything?



# Our First “Real” Query

```
SELECT f.DRINKER
FROM FREQUENTS AS f, SERVES AS s
WHERE f.BAR = s.BAR AND s.BEER = "SSTP"
```

FREQUENTS

```
('Luis', 'Bar1')
('Luis', 'Bar2')
```

SERVES

```
('Bar1', 'SSTP')
('Bar1', 'Modelo')
('Bar2', 'Blue_Moon' )
('Bar2', 'SSTP' )
```

OUTPUT

```
('Luis')
('Luis')
```

# Our First “Real” Query

LIKES (DRINKER, BEER)

FREQUENTS (DRINKER, BAR)

SERVES (BAR, BEER)

Who goes to a bar serving Sam Smith Taddy Porter? (‘SSTP’)

```
SELECT DISTINCT f.DRINKER  
FROM FREQUENTS AS f, SERVES AS s  
WHERE f.BAR = s.BAR AND s.BEER = "SSTP"
```

Are we missing anything?

- The DISTINCT keyword

# Our First “Real” Query

LIKES (DRINKER, BEER)

FREQUENTS (DRINKER, BAR)

SERVES (BAR, BEER)

Who goes to a bar serving Sam Smith Taddy Porter? (‘SSTP’)

```
SELECT DISTINCT f.DRINKER  
FROM FREQUENTS AS f, SERVES AS s  
WHERE f.BAR = s.BAR AND s.BEER = "SSTP"
```

Closely related to RC! Same as:

▷  $\{f.DRINKER | FREQUENTS(f) \wedge \exists(s)(SERVES(s) \wedge f.BAR = s.BAR \wedge s.BEER = \text{'SSTP'})\}$

# AS Keyword

```
SELECT DISTINCT f.DRINKER  
FROM FREQUENTS AS f, SERVES AS s  
WHERE f.BAR = s.BAR AND s.BEER = "SSTP"
```

What does AS do?

- ▷ Rename ( $\rho$ ) from Relational Algebra!
- ▷ Works on tables as well as attributes
- ▷ Actual keyword is optional
- ▷ Why bother? To create a more meaningful name

```
SELECT DISTINCT f.DRINKER "SSTP_Drinkers"  
FROM FREQUENTS f, SERVES s  
WHERE f.BAR = s.BAR AND s.BEER = "SSTP"
```

## Select Clause

Attribute	Example
Attribute list	$d.attr1, d.attr2$
All attributes	*
$\langle \text{table name} \rangle . *$	FREQUENTS.*
$\langle \text{alias name} \rangle . *$	$f . *$
$\langle \text{math equation} \rangle$	$1 + 3$
$\langle \text{constant} \rangle$	'CPA', 3
$\langle \text{function} \rangle$	NOW, CONCAT, COALESCE
Eliminate duplicates	DISTINCT

# Where Clause

- `<attribute> = <value>`
- `<attribute> BETWEEN [value1] AND [value2]`
- `<attribute> IN ([value1], [value2], ...)`
- `<attribute> LIKE 'SST%'`
- `<attribute> LIKE 'SST_'`
- `<attribute> IS NULL` and `[attribute] IS NOT NULL`
- Logical combinations with AND and OR
- Mathematical functions `<>`, `!=`, `>`, `<`, ...
- Subqueries ...

# Subqueries

Can have a subquery in the WHERE clause

Linked with keywords

- EXISTS
- IN
- ALL
- SOME

# Subqueries

Can have a subquery in the WHERE clause

Linked with keywords

- EXISTS
  - ▷ EXISTS <subquery>
  - ▷ If the subquery returns at least one tuple, the clause evaluates to TRUE
  - ▷ NOT EXISTS?
- IN
- ALL
- SOME



# Subqueries

Can have a subquery in the WHERE clause

Linked with keywords

- EXISTS

- IN

  - ▷ `<expression> IN <subquery> / <expression> NOT IN <subquery>`

  - ▷ How does IN work?

- ALL

- SOME

What is an expression in this context?

# Subqueries

Can have a subquery in the WHERE clause

Linked with keywords

- EXISTS
- IN
- ALL
  - ▷ `<expression> <boolOP> ALL <subquery>`
  - ▷ TRUE if every item in the subquery makes the boolOp evaluate to TRUE
- SOME
  - ▷ `<expression> <boolOP> SOME/ANY <subquery>`
  - ▷ TRUE if some item in the subquery can make the boolOp evaluate to TRUE

# Subqueries

How do subqueries work?

- As we iterate over the tuples of the outer query, the inner query is evaluated for each tuple.
- Some can be evaluated just once
  - ▷ E.g., a subquery that returns the number of BARS that are frequented
- Some require the subquery to be evaluated for every value assignment in the outer query
  - ▷ E.g., a subquery that returns the number of BARS that each DRINKER goes to
  - ▷ Correlated subqueries

# Subquery Example #1: IN

LIKES(DRINKER, BEER)

Q: Who likes 'PBR' and 'Corona'?

1. Figure out who likes 'PBR'
2. Use the subquery to make sure they also like 'Corona'

# Subquery Example #1: IN

LIKES(DRINKER, BEER)

Q: Who likes 'PBR' and 'Corona'?

1. Figure out who likes 'PBR'
2. Use the subquery to make sure they also like 'Corona'

```
SELECT DISTINCT l.DRINKER  
FROM LIKES l  
WHERE l.BEER = 'PBR'
```

# Subquery Example #1: IN

LIKES(DRINKER, BEER)

Q: Who likes 'PBR' and 'Corona'?

```
SELECT DISTINCT l.DRINKER  
FROM LIKES l  
WHERE l.BEER = 'PBR'  
         AND l.DRINKER IN (people who like Corona)
```

# Subquery Example #1: IN

LIKES(DRINKER, BEER)

Q: Who likes 'PBR' and 'Corona'?

```
SELECT DISTINCT 1.DRINKER
FROM LIKES 1
WHERE 1.BEER = 'PBR'
        AND 1.DRINKER IN (
                SELECT 12.DRINKER
                FROM LIKES 12
                WHERE 12.BEER = 'Corona')
```

What is the subquery returning?

# Subquery Example #1: IN

Q: Who likes 'PBR' and 'Corona'?

Many subqueries can be written as JOINS

▷ People find it easier to reason about it one way or the other

```
SELECT DISTINCT 1.DRINKER
FROM LIKES 1
WHERE 1.BEER = 'PBR'
        AND 1.DRINKER IN (
                SELECT 12.DRINKER
                FROM LIKES 12
                WHERE 12.BEER = 'Corona')
```

```
SELECT DISTINCT 11.DRINKER
FROM LIKES 11, LIKES 12
WHERE 11.DRINKER = 12.DRINKER
        AND 11.BEER = 'PBR'
        AND 12.BEER = 'Corona'
```



## Subquery Example #2: SOME

RATES (DRINKER, BEER, SCORE)

Q: List the beers that are not Luis' favorite.

What does it mean, in terms of RATES, when we say favorite?

## Subquery Example #2: SOME

RATES (DRINKER, BEER, SCORE)

Q: List the beers that are not Luis' favorite.

1. Find the beers that Luis likes
2. Use the subquery to select every non-favorite beer

```
SELECT r.BEER  
FROM RATES r  
WHERE r.DRINKER = 'Luis'
```

## Subquery Example #2: SOME

RATES (DRINKER, BEER, SCORE)

Q: List the beers that are not Luis' favorite.

What does it mean, in terms of RATES, when we say favorite?

```
SELECT r.BEER  
FROM RATES r  
WHERE r.DRINKER = 'Luis' AND (this beer score must not be the highest)
```

## Subquery Example #2: SOME

RATES (DRINKER, BEER, SCORE)

Q: List the beers that are not Luis' favorite.

What does it mean, in terms of RATES, when we say favorite?

```
SELECT r.BEER
FROM RATES r
WHERE r.DRINKER = 'Luis' AND r.SCORE < SOME (
    SELECT r2.SCORE
    FROM RATES r2
    WHERE r2.DRINKER = 'Luis')
```

What is the subquery returning?

# Views

Q: List the beers that are not Luis' favorite.

“Common” (non-materialized) views are just macros

- Unexecuted query
- Can be used in place of a table
- Convenient way to simplify a query
- Query is executed when view is used by another query
- Its results are not stored

```
CREATE VIEW LUIS_BEERS AS  
SELECT *  
FROM RATES r  
WHERE r.DRINKER = 'Luis'
```

# Views

Q: List the beers that are not Luis' favorite.

```
CREATE VIEW LUIS_BEERS AS  
SELECT *  
FROM RATES r  
WHERE r.DRINKER = 'Luis'
```

```
SELECT r.BEER  
FROM LUIS_BEERS  
WHERE r.SCORE < SOME (  
    SELECT r2.SCORE  
    FROM LUIS_BEERS r2)
```

## Subquery Example #3: NOT EXISTS

LIKES (DRINKER, BEER)

FREQUENTS (DRINKER, BAR)

SERVES (BAR, BEER)

Q: Who likes all of the beers that Luis likes?

## Subquery Example #3: NOT EXISTS

LIKES (DRINKER, BEER)

FREQUENTS (DRINKER, BAR)

SERVES (BAR, BEER)

Q: Who likes all of the beers that Luis likes?

- There does not exist a beer that Luis likes that is not also liked by these drinkers
- Every beer Luis likes is liked by these drinkers BUUUUT they might like other beers as well



## Subquery Example #3: NOT EXISTS

LIKES (DRINKER, BEER)

FREQUENTS (DRINKER, BAR)

SERVES (BAR, BEER)

Q: Who likes all of the beers that Luis likes?

```
SELECT l.DRINKER
FROM LIKES l
WHERE NOT EXISTS (a beer Luis likes that is not
    also liked by l.DRINKER)
```

## Subquery Example #3: NOT EXISTS

LIKES (DRINKER, BEER)

FREQUENTS (DRINKER, BAR)

SERVES (BAR, BEER)

Q: Who likes all of the beers that Luis likes?

1. Beer that Luis likes

```
SELECT l2.BEER  
FROM LIKES l2  
WHERE l2.DRINKER = 'Luis'
```

## Subquery Example #3: NOT EXISTS

LIKES (DRINKER, BEER)

FREQUENTS (DRINKER, BAR)

SERVES (BAR, BEER)

Q: Who likes all of the beers that Luis likes?

```
SELECT 1.DRINKER
FROM LIKES 1
WHERE NOT EXISTS (
    SELECT 12.BEER
    FROM LIKES 12
    WHERE 12.DRINKER = 'Luis' AND 12.BEER NOT IN (
        the set of beers liked by 1.DRINKER) )
```

## Subquery Example #3: NOT EXISTS

LIKES (DRINKER, BEER)

FREQUENTS (DRINKER, BAR)

SERVES (BAR, BEER)

Q: Who likes all of the beers that Luis likes?

2. Beer that 1.DRINKER likes

```
SELECT 13.BEER  
FROM LIKES 13  
WHERE 13.DRINKER = 1.DRINKER
```

## Subquery Example #3: NOT EXISTS

Q: Who likes all of the beers that Luis likes?

Putting it all together

```
SELECT 1.DRINKER
FROM LIKES 1
WHERE NOT EXISTS (
    SELECT 12.BEER
    FROM LIKES 12
    WHERE 12.DRINKER = 'Luis' AND 12.BEER NOT IN (
        SELECT 13.beer
        FROM LIKES 13
        WHERE 13.DRINKER = 1.DRINKER) )
```

## Subquery Example #3: NOT EXISTS

Q: Who likes all of the beers that Luis likes?

Putting it all together

```
SELECT l.DRINKER
FROM LIKES l
WHERE NOT EXISTS (
    SELECT l2.BEER
    FROM LIKES l2
    WHERE l2.DRINKER = 'Luis' AND l2.BEER NOT IN (
        SELECT l3.beer
        FROM LIKES l3
        WHERE l3.DRINKER = l.DRINKER) )
```

Same as:

▷  $\{l.DRINKER | LIKES(l) \wedge \neg \exists (l_2)(LIKES(l_2) \wedge l_2.DRINKER = 'Luis' \wedge \neg \exists (l_3)(LIKES(l_3) \wedge l_3.DRINKER = l.DRINKER \wedge l_3.BEER = l_2.BEER))\}$

# Some Closing Notes

## Style

- ▷ Declarative SQL codes tend to be very short
- ▷ Good because effort, bugs  $\propto$  code length
- ▷ Bad because sometimes difficult to understand!

# Some Closing Notes

## Style

- ▷ Declarative SQL codes tend to be very short
- ▷ Good because effort, bugs  $\propto$  code length
- ▷ Bad because sometimes difficult to understand!

Hence, style is important. Some suggestions

- ▷ Always alias tuple variables
- ▷ Always indent carefully
- ▷ Only one major keyword per line (SELECT, FROM, etc.)
- ▷ Pick a capitalization schema and religiously stick to it
- ▷ Make frequent use of views



# Questions?